METRUM
RESEARCH GROUP

Torsten

A Prototype Library for Bayesian Pharmacometrics Modeling in Stan

User Manual

Charles Margossian and Bill Gillespie

Torsten Version 0.82
for Stan Version 2.14.0

Tuesday January 31st 2017

CONTENTS

ACKNOWLEDGEMENTS

1. INTRODUCTION

1.1. **Preface.**

Stan is an open source probabilistic language designed primarily to do Bayesian data analysis [1]. Several of its features make it a powerful tool to specify and fit complex models. Notably, its language is extremely flexible and its No U-Turn Sampler (NUTS), an adaptive Hamiltonian Monte Carlo algorithm, has proven more efficient than commonly used Monte Carlo Markov Chains (MCMC) samplers for complex high dimensional problems [2]. Our goal is to harness these innovative features and make Stan a better software for pharmacometrics modeling. Our efforts are twofold:

(1) We contribute to the development of new mathematical tools, such as functions that support differential equations based models, and implement them directly into Stan's core language.
(2) We develop Torsten, an extension with specialized pharmacometrics functions.

Throughout the process, we have been working very closely with Stan's development team. We have benefited immensely from their mentoring, advice, and feedback. Just like Stan, Torsten is an open source project that fosters collaborative work. Interested in contributing? Shoot us an e-mail and we will help you help us (`charlesm@metrumrg.com` and `billg@metrumrg.com`)!

Torsten is licensed under the BSD 3-clause license.

> **WARNING:** The current version of Torsten is a *prototype*. It is being released for review and comment, and to support limited research applications. It has not been rigorously tested and should not be used for critical applications without further testing or cross-checking by comparison with other methods.
>
> We encourage interested users to try Torsten out and are happy to assist. Please report issues, bugs, and feature requests on our GitHub page: `https://github.com/charlesm93/stan`.

1.2. **Installing Torsten.**

Installation files are available on GitHub: `https://github.com/charlesm93/example-models/tree/torsten-0.82/PKPD/torsten`

There is currently no mechanism to install Torsten on top of your version of Stan. This is still a work in progress. In the meantime, we offer a version of Stan with Torsten built inside of it. Torsten 0.82 works with Stan 2.14.0. Torsten is built inside the Stan and Stan-math repositories and is agnostic to the interface the users use. We offer support for installing Stan with CmdStan[1].

The easiest way to install the CmdStan interface with Stan and Torsten is to run the bash file `setupTorsten.sh`. You can also do it manually on a terminal window:

Download CmdStan:
```
git clone https://github.com/stan-dev/cmdstan.git
cd cmdstan
git checkout release/v2.14.0
```

---

[1]If you need assistance to set up Torsten with another interface, shoot me an e-mail at charlesm@metrumrg.com. I'm happy to help.

Download Stan with Torsten:
```
git clone https://github.com/charlesm93/stan.git
rm -r stan_2.14.0
mv stan stan_2.14.0
cd stan_2.14.0
git checkout torsten-master
```

Download Stan_math with Torsten: `cd lib`
```
git clone https://github.com/charlesm93/math.git
mv math stan_math_2.14.0
cd stan_math_2.14.0
git checkout torsten-master
```

1.3. **Overview.**

Torsten is a prototype Pharmacokinetic/Pharmacodynamic (PKPD) model library for use in Stan 2.14.0. The current version includes:

- Specific linear compartmental models:
  - One compartment model with first order absorption
  - Two compartment model with elimination from and first order absorption into central compartment
- General linear compartmental model described by a system of first-order linear Ordinary Differential Equations (ODEs).
- General compartmental model described by a system of first order ODEs

The models and data format are based on NONMEM®[2]/NMTRAN/PREDPP conventions including:

- Recursive calculation of model predictions
  - This permits piecewise constant covariate values
- Bolus or constant rate inputs into any compartment
- Handles single dose and multiple dose histories
- Handles steady-state dosing histories for specific and general <u>linear</u> models
- Implemented NMTRAN data items include: TIME, EVID, CMT, AMT, RATE, ADDL, II, SS

This library provides Stan language functions that calculate amounts in each compartment, given an event schedule and an ODE system.

1.4. **Implementation details.**

- Stan version 2.14.0 (`http://mc-stan.org/`)
- All functions are programmed in C++ and are compatible with the Stan-math automatic differentiation library [3]
- All functions can be called directly in a Stan file in a manner identical to other built-in functions
- One and two compartment models: hand-coded analytical solutions
- General linear compartment models with semi-analytical solutions using a built-in Matrix Exponential function (Pade approximation coupled with scaling and squaring [4])

---

[2]NONMEM® is licensed and distributed by ICON Development Solutions.

- General compartment models with numerical solutions to ODEs using built-in ODE integrators in Stan (Runge-Kutta 4th/5th and backward differentiation methods from CVODES library), with adjustable tuning parameters

1.5. **Development plans.**

Our current plans for future development of Torsten include the following:

- A system to easily share packages of Stan functions (written in C++ or in the Stan language)
- Steady state calculation for the General compartment model
    - This requires a solver for nonlinear algebraic equations (aka a root solver)
- Optimize Matrix exponential functions
    - Function for the action of Matrix Exponential on a vector
    - Hand-coded gradients
    - Special algorithm for matrices with special properties
- Fix issue that arises when computing the adjoint of the lag time parameter (in a dosing compartment) evaluated at $t_{lag} = 0$.
- Make the following arguments optional
    - `biovar` and `tlag`, respectively used for the bioavailability fraction and the lag times in each compartment
    - Tuning parameters of the ODE integrators for the general compartmental function.
- Extend formal tests
    - We want more C++ Google unit tests to address cases users may encounter
    - Comparison with simulations from the R package *mrgsolve* and the software NONMEM®
    - Recruit non-developer users to conduct beta testing

1.6. **Updates since Torsten 0.81.**

- Torsten is now up to date with Stan version 2.14.0
- We split the parameter argument into three arguments: `pMatrix` (parameters for ODEs), `biovar` (parameters for the bioavailability), and `tlag` (parameters for lag times). This gives the users more control over which arguments get passed as data and as parameters. The model is most efficient when all fixed values are passed as data.
- We changed the order of the arguments: the user first passes the data arguments and then the parameter arguments. This is because in future versions we will want optional arguments to be passed last.
- The user can now pass parameter arguments as 1D or 2D arrays, depending on whether the parameters are constant or change from one event to the other.
- We have significantly increased the number of unit tests – but we still have more to do!
- The unit tests now check automatic differentiation against finite differentiation calculations.
- Fixed minor bugs and report issues.

2. Using Torsten

The reader should have a basic understanding of how Stan works before tackling this chapter. There are excellent resources online to get started with Stan (`http://mc-stan.org/documentation/`).

In this section we go through the different functions Torsten adds to Stan. It will be helpful to apply these functions to a simple example. We have uploaded code and data on `https://github.com/charlesm93/example-models/tree/torsten-0.82/PKPD/torsten`.

2.1. Example 1: Two Compartment Model.

We model drug absorption in a single patient and simulate plasma drug concentrations:

- Single Dose: 5000 mg
- PK: plasma concentrations of parent drug (c)
- PK measured at 0, 0.25, 0.50, 0.75, 1, 1.25, 1.50, 1.75, 2, 4, 6, 8, 10, 12, 14, 16, 18, and 20 hours after dosing

The plasma concentration (c) are simulated according to the following equations:

$$
\begin{aligned}
\log(c) &\sim N\left(\log(\widehat{c}), \sigma^2\right) \\
\widehat{c} &= f_{2cpt}(t, CL, Q, V_2, V_3, k_a) \\
(CL, Q, V_2, V_3, ka) &= \left(5\ \text{L/h}, 8\ \text{L/h}, 20\ \text{L}, 70\ \text{L}, 1.2\ \text{h}^{-1}\right) \\
\sigma &= 0.05
\end{aligned}
$$

The data are generated using the R package *mrgsolve*[3], see `TwoCptModelSimulation.R`. We show the results obtained when using the function `PKModelTwoCpt`, which computes solutions to the ODEs analytically.

2.2. Linear One and Two Compartment Model Function.

The one and two compartment model functions have the form:

```
<model name>(time, amt, rate, ii, evid, cmt, addl, ss,
             theta, biovar, tlag)
```

There is no need to skip a line, but we do so to separate what we call *event* arguments and *model* arguments. The event arguments describe the event schedule of the clinical trial. `time`, `amt`, `rate`, `ii` are arrays of real and `evid`, `cmt`, `addl`, `ss` arrays of integers. All arrays have the same length, which corresponds to the number of events.

Next we have the *model* arguments: `theta` contains the ODE parameters, `biovar` the bioavailability fraction in each compartment (sometimes denoted as F and called *biovariability*), and `tlag` the lag times in each compartment. The model arguments are 2-dimensional arrays. The first dimension provides a container for the parameters and the length of that container is simply the number of parameters. The second dimension allows for variations in the parameters from one event to the other. The length of the 2D container is the number of events. The *ith* row contains an array of parameters for the time interval

---

[3]https://github.com/metrumresearchgroup/mrgsolve

`[time[i-1], time[i]]`. Often times, the parameters are constant for all events and the user can then pass a 2D array of length 1 or, even better, a 1D array.

The options for *model name* are:

- PKModelOneCpt
- PKModelTwoCpt

which respectively correspond to the one and two compartment model with first order absorption (figure 1). A vector in `theta` is expected to contain parameters CL, $V2$, and ka for the one compartment case, and CL, Q, $V2$, $V3$, and ka for the two compartments case, <u>in this order</u>. Setting ka to 0 eliminates the fist-order absoprtion. `biovar` contains the bioavailability fraction of each compartment (non-effective if set to 1) and `tlag` the lag time in each compartment (non-effective if set to 0).
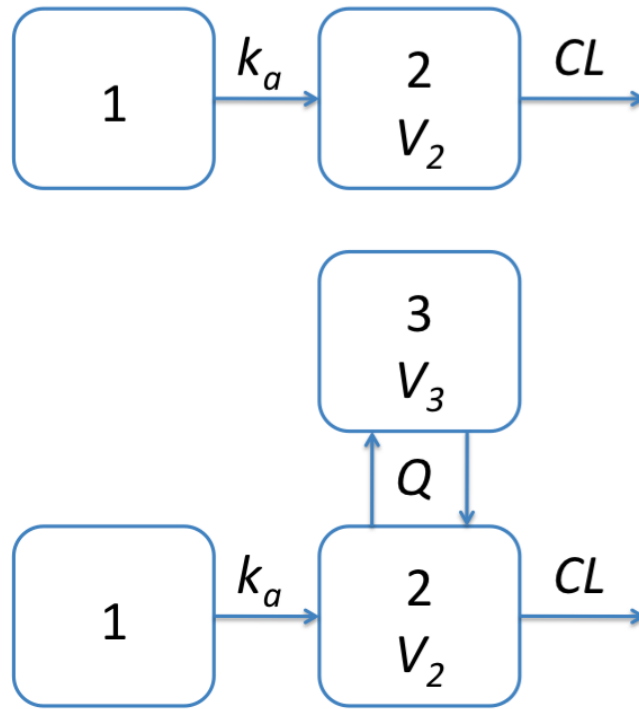

FIGURE 1. One and two compartment models with first order absorption implemented in Torsten.

`PKModelTwoCpt` can be used to fit example 1, see `TwoCptModel.stan`. We are interested in evaluating the ODE parameters, stored in `theta`. The bioavailability fraction and the lag times on the other hand are fixed, and we therefore declare `biovar` and `tlag` in the transformed data block. Four MCMC chains of 2000 iterations were simulated. The first 1000 iteration of each chain were discarded. Thus 1000 MCMC samples were used for the subsequent analyses.

Result. The MCMC history plots (figure 3) suggest that the 4 chains have converged to common distributions for all of the key model parameters. The fit to the plasma concentration data (figure 5) are in close agreement with the data, which is not surprising since the fitted model is identical to the one used to simulate the data. Similarly the parameter estimates summarized in Table 1 are consistent with the values used for simulation.

FIGURE 2. Stan language for fitting a two compartment model using the `PKModelTwoCpt` function (abstract)

```
data {
  int<lower = 1> nt; # number of events
  int<lower = 1> nObs; # number of observation
  int<lower = 1> iObs[nObs]; # index of observation
  int<lower = 1> cmt[nt];
  int evid[nt];
  int addl[nt];
  int ss[nt];
  real amt[nt];
  real time[nt];
  real rate[nt];
  real ii[nt];

  vector<lower = 0>[nObs] cObs; # observed concentration (Dependent Variable)
}

transformed data {
                        ⋮

  biovar[1] = 1;
  biovar[2] = 1;
  biovar[3] = 1;

  tlag[1] = 0;
  tlag[2] = 0;
  tlag[3] = 0;


}
                        ⋮
parameters {
  real<lower = 0> CL;
  real<lower = 0> Q;
  real<lower = 0> V2;
  real<lower = 0> V3;
  real<lower = 0> ka;
  real<lower = 0> sigma;
}

transformed parameters {
                        ⋮

  theta[1] = CL;
  theta[2] = Q;
  theta[3] = V2;
  theta[4] = V3;
  theta[5] = ka;

  x = PKModelTwoCpt(time, amt, rate, ii, evid, cmt, addl, ss,
                    theta, biovar, tlag);

  cHat = col(x, 2) ./ V2; # get concentration in the central compartment

  for(i in 1:nObs){
    cHatObs[i] = cHat[iObs[i]];# predictions for observed data records
  }
}
                        ⋮
```
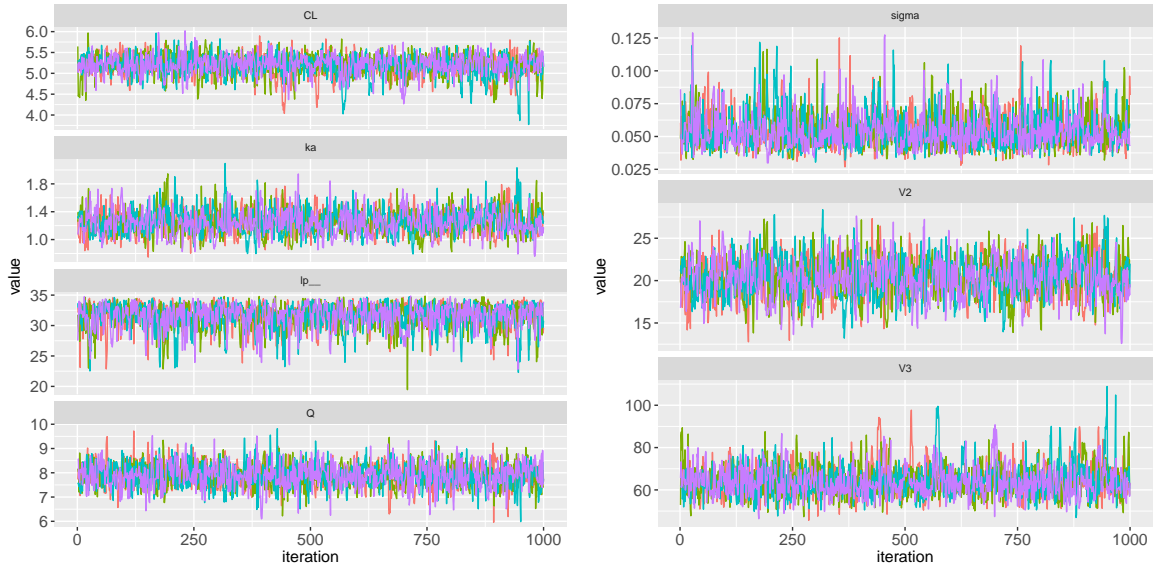
FIGURE 3. MCMC history plots for the parameters of a two compartment model with first order absorption (each color corresponds to a different chain)
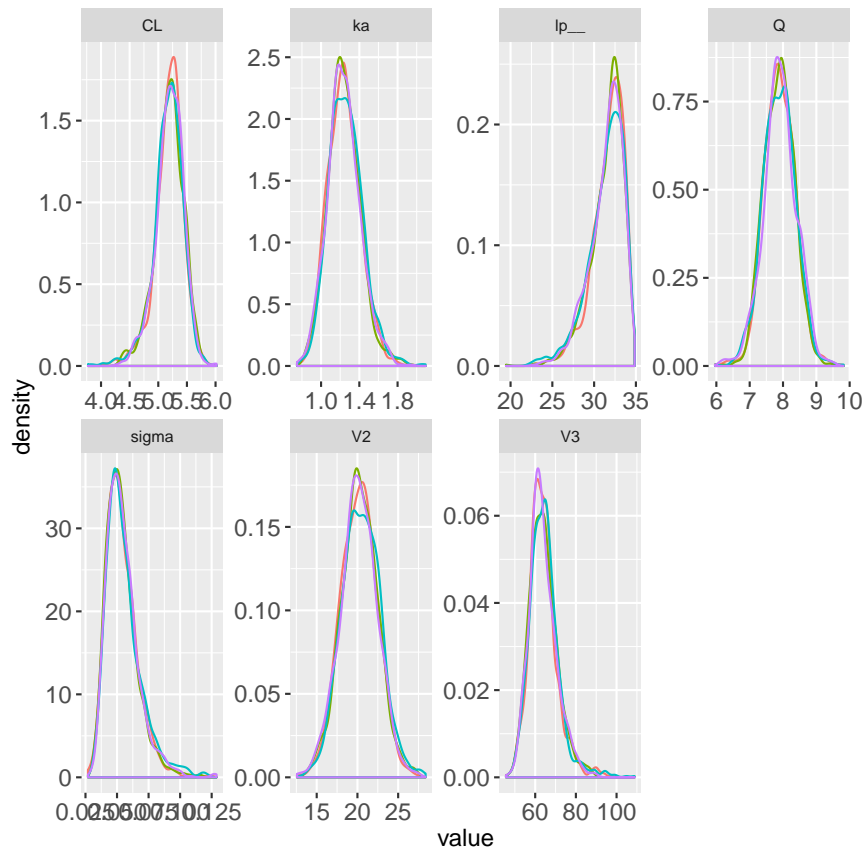

FIGURE 4. Posterior Marginal Densities of the Model Parameters of a two compartment model with first order absorption (each color corresponds to a different chain)

FIGURE 5. Predicted (posterior median and 90 % credible intervals) and observed plasma drug concentrations of a two compartment model with first order absorption

TABLE 1. Summary of the MCMC simulations of the marginal posterior distributions of the model parameters

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
CL	5.19	0.01	0.25	4.60	5.05	5.21	5.35	5.62	971.04	1.00
Q	7.90	0.01	0.47	6.97	7.58	7.93	8.21	8.79	1662.02	1.00
V2	20.25	0.08	2.27	15.71	18.74	20.23	21.76	24.67	896.67	1.00
V3	64.10	0.22	6.98	52.60	59.48	63.38	67.74	80.65	1032.66	1.00
ka	1.25	0.00	0.17	0.94	1.13	1.24	1.35	1.62	940.99	1.00

2.3. General Linear ODE Model Function.

A general linear ODE model refers to a model that may be described in terms of a system of first order linear differential equations with (piecewise) constant coefficients, i.e., a differential equation of the form:

$$y'(t) = Ky(t)$$

where K is a matrix. For example K for a two compartment model with first order absorption is:

$$K = \begin{bmatrix} -k_a & 0 & 0 \\ k_a & -(k_{10} + k_{12}) & k_{21} \\ 0 & k_{12} & -k_{21} \end{bmatrix}$$

where $k_{10} = CL/V2$, $k_{12} = Q/V2$, and $k_{21} = Q/V2$.

The linear ODE model function has the form:

```
linOdeModel(time, amt, rate, ii, evid, cmt, addl, ss,
            system, biovar, tlag)
```

where `system` is an array of constant rate matrices. The length of the array is the number of events. If the rate matrix is constant for all events, the user may pass an array of length 1 or directly the matrix `K`. `system` contains all the ODE parameters, so we no longer need `theta`.

FIGURE 6. Stan language for fitting a two compartment model using the `linOdeModel` function (abstract)

```
transformed parameters {
  matrix[3, 3] K;
  real k10 = CL / V2;
  real k12 = Q / V2;
  real k21 = Q / V3;
  vector<lower = 0>[nTheta] theta[1];
  vector<lower = 0>[nt] cHat;
  vector<lower = 0>[nObs] cHatObs;
  matrix<lower = 0>[nt, 3] x;

  K = rep_matrix(0, 3, 3);

  K[1, 1] = -ka;
  K[2, 1] = ka;
  K[2, 2] = -(k10 + k12);
  K[2, 3] = k21;
  K[3, 2] = k12;
  K[3, 3] = -k21;

  x = linOdeModel(time, amt, rate, ii, evid, cmt, addl, ss,
                  K, biovar, tlag);

  cHat = col(x, 2) ./ V1;

  for(i in 1:nObs){
    cHatObs[i] = cHat[iObs[i]]; # predictions for observed data records
  }
}

model{
  logCObs ~ normal(log(cHatObs), sigma);
}
```

2.4. General ODE Model Function.

Torsten may be used to fit models described by a system of first-order ODEs, i.e., differential equations of the form:

$$y'(t) = f(t, y(t))$$

where y and f are vector-valued functions.

The general ODE model functions have the form:

```
<model_name>(ODE_system, nCmt,
             time, amt, rate, ii, evid, cmt, addl, ss,
             theta, biovar, tlag,
             rel_tol, abs_tol, max_step)
```

where `ODE_system` is a system of first-order ODEs defined in the function block of Stan (see section 19.2 of the Stan reference manual) and `nCmt` is the number of compartments (or, equivalently, the number of ODEs) in the model. `rel_tol`, `abs_tol`, and `max_step` are the tuning parameters for the ODE integrator: respectively the relative tolerance, the absolute tolerance, and the maximum number of steps. Which value to pick for the tuning parameters depends on the specifics of the ODEs. Reducing the tolerance parameters and increasing the number of steps make for a more robust integrator but can significantly slow down the algorithm. We here recommend default values, but users should be prepared to adjust these parameters. As a starting point, we suggest: `rel_tol` = 1e-8, `abs_tol` = 1e-8 and `max_step` = 1e+8. For more details see Stan's reference manual.

The options for `model_name` are:

- generalOdeModel_rk45
- generalOdeModel_bdf

They respectively call the built-in Runge-Kutta 4th/5th integrator, recommended for non-stiff ODEs, and BDF integrator, recommended for stiff ODEs.

TABLE 2. Summary: Arguments of Torsten functions.

model	function name	argument names	parameters in `theta`
one compartment model with first order absorption	`PKModelOneCpt`	`time, amt, rate, ii, evid, cmt, addl, ss, theta, biovar, tlag`	CL, V_2, k_a
two compartment model with first order absorption	`PKModelTwoCpt`	`time, amt, rate, ii, evid, cmt, addl, ss, theta, biovar, tlag`	CL, Q, V_2, V_3, k_a
general linear compartment model	`linOdeModel`	`time, amt, rate, ii, evid, cmt, addl, ss, system, biovar, tlag`	NA: pass in constant rate matrix instead of `theta`
general compartment models	`genOdeModel_*`	`ODE_system, nCmt, time, amt, rate, ii, evid, cmt, addl, ss, theta, biovar, tlag, rel_tol, abs_tol, max_num_steps`	Parameters that get passed to ODE system

Figure 7. Stan language for fitting a two compartment model using the `genOdeModel_rk45` function (abstract)

```
functions{
  # define ODE system for two compartment model
  real[] twoCptModelODE(real t,
                        real[] x,
                        real[] theta,
                        real[] dummy_real,
                        int[] dummy_int){
    real Q = theta[1];
    real CL = theta[2];
    real V2 = theta[3];
    real V3 = theta[4];
    real ka = theta[5];
    real k12 = Q / V2;
    real k21 = Q / V3;
    real k10 = CL / V2;
    real y[3];

    y[1] = -ka * x[1];
    y[2] = ka * x[1] - (k10 + k12)*x[2] + k21*x[3];
    y[3] = k12 * x[2] - k21 * x[3];

    return y;
  }
}
                          .
                          .
                          .
transformed parameters {

                          .
                          .
                          .
  theta[1] = CL;
  theta[2] = Q;
  theta[3] = V1;
  theta[4] = V2;
  theta[5] = ka;

  x = generalCptModel_rk45(twoCptModelODE, 3,
                           time, amt, rate, ii, evid, cmt, addl, ss,
                           theta, biovar, tlag,
                           1e-8, 1e-8, 1e8);
                           .
                           .
                           .
```

3. ADDITIONAL EXAMPLES

Code for examples can be found on GitHub: `https://github.com/charlesm93/example-models/tree/torsten-0.82/PKPD/torsten`.

All the files to run a model are stored under the directory that bears the model's name. There are four files per example:

- `<model name>.stan`
- `<model name>.data.R`
- `<model name>.init.R`
- `<model name>Simulation.R`

`data.R` contains the data we fit the model to and `init.R` the initial estimates of the parameters. These two files are generated using `Simulation.R`. The `R` folder contains R scripts to compile and run the models, as well as code to output diagnostic plots and statistics.

3.1. **Effect Compartment Model.**

Let us expand example 1 to a population model fitted to the combined data from phase I and phase IIa studies. The parameters exhibit inter-individual variations (IIV), due to both random effects and to the patients' body weight, treated as a covariate and denoted bw:

Population Model for Plasma Drug Concentration (c).

$$
\begin{aligned}
\log\left(c_{ij}\right) &\sim N\left(\log\left(\widehat{c}_{ij}\right), \sigma^2\right) \\
\widehat{c}_{ij} &= f_{2cpt}\left(t_{ij}, D_j, \tau_j, CL_j, Q_j, V_{1j}, V_{2j}, k_{aj}\right) \\
\log\left(CL_j, Q_j, V_{ssj}, k_{aj}\right) &\sim N\left(\log\left(\widehat{CL}\left(\frac{bw_j}{70}\right)^{0.75}, \widehat{Q}\left(\frac{bw_j}{70}\right)^{0.75}, \widehat{V}_{ss}\left(\frac{bw_j}{70}\right), \widehat{k}_a\right), \Omega\right) \\
V_{1j} &= f_{V_1} V_{ssj} \qquad V_{2j} = \left(1 - f_{V_1}\right) V_{ssj} \\
\left(\widehat{CL}, \widehat{Q}, \widehat{V}_{ss}, \widehat{k}_a, f_{V_1}\right) &= \left(10\ \mathrm{L/h}, 15\ \mathrm{L/h}, 140\ \mathrm{L}, 2\ \mathrm{h}^{-1}, 0.25\right) \\
\Omega &= \begin{pmatrix} 0.25^2 & 0 & 0 & 0 \\ 0 & 0.25^2 & 0 & 0 \\ 0 & 0 & 0.25^2 & 0 \\ 0 & 0 & 0 & 0.25^2 \end{pmatrix}, \quad \sigma = 0.1
\end{aligned}
$$

Furthermore we add a fourth compartment in which we measure a PD effect (figure 8).

FIGURE 8. Effect Compartment Model

Effect Compartment Model for PD response (R).

$$
\begin{aligned}
R_{ij} &\sim N\left(\widehat{R}_{ij}, \sigma_R^2\right) \\
\widehat{R}_{ij} &= \frac{E_{max} c_{eij}}{EC_{50j} + c_{eij}} \\
c'_{e\cdot j} &= k_{e0j}\left(c_{\cdot j} - c_{e\cdot j}\right) \\
\log\left(EC_{50j}, k_{e0j}\right) &\sim N\left(\log\left(\widehat{EC}_{50}, \widehat{k}_{e0}\right), \Omega_R\right) \\
\left(E_{max}, \widehat{EC}_{50}, \widehat{k}_{e0}\right) &= (100, 100.7, 1) \\
\Omega_R &= \begin{pmatrix} 0.2^2 & 0 \\ 0 & 0.25^2 \end{pmatrix}, \quad \sigma_R = 10
\end{aligned}
$$

The PK and the PD data are simulated using the following treatment.

- Phase I study
    - Single dose and multiple doses
    - Parallel dose escalation design
    - 25 subjects per dose
    - Single doses: 1.25, 5, 10, 20, and 40 mg
    - PK: plasma concentration of parent drug (c)
    - PD response: Emax function of effect compartment concentration (R)
    - PK and PD measured at 0.083, 0.167, 0.25, 0.5, 0.75, 1, 2, 3, 4, 6, 8, 12, 18, and 24 hours
- Phase IIa trial in patients
    - 100 subjects
    - Multiple doses: 20 mg
    - sparse PK and PD data (3-6 samples per patient)

FIGURE 9. Stan language for fitting an effect compartment model using `linOdeModel` (abstract)

```
transformed parameters {
  for(j in 1:nSubjects){
                    .
                    .
                    .
  Omega = quad_form_diag(rho, omega);

  for(j in 1:nSubjects){
    CL[j] = exp(logtheta[j, 1]) * (weight[j] / 70)^0.75;
    Q[j] = exp(logtheta[j, 2]) * (weight[j] / 70)^0.75;
    V1[j] = exp(logtheta[j, 3]) * weight[j] / 70;
    V2[j] = exp(logtheta[j, 4]) * weight[j] / 70;
    ka[j] = exp(logtheta[j, 5]);
    ke0[j] = exp(logKe0[j]);
    EC50[j] = exp(logEC50[j]);

    k10 = CL[j] / V1[j];
    k12 = Q[j] / V1[j];
    k21 = Q[j] / V2[j];

    K = rep_matrix(0, 4, 4);
    K[1, 1] = -ka[j];
    K[2, 1] = ka[j];
    K[2, 2] = -(k10 + k12);
    K[2, 3] = k21;
    K[3, 2] = k12;
    K[3, 3] = -k21;
    K[4, 2] = ke0[j];
    K[4, 4] = -ke0[j];

    ke0[j] = exp(logKe0[j]);
    EC50[j] = exp(logEC50[j]);

    K = rep_matrix(0, 4, 4);

    K[1, 1] = -ka[j];
    K[2, 1] = ka[j];
    K[2, 2] = -(k10 + k12);
    K[2, 3] = k21;
    K[3, 2] = k12;
    K[3, 3] = -k21;
    K[4, 2] = ke0[j];
    K[4, 4] = -ke0[j];

    x[start[j]:end[j],] = linOdeModel(time[start[j]:end[j]],
                              amt[start[j]:end[j]], rate[start[j]:end[j]],
                              ii[start[j]:end[j]], evid[start[j]:end[j]],
                              cmt[start[j]:end[j]], addl[start[j]:end[j]],
                              ss[start[j]:end[j]], K, biovar, tlag);

    cHat[start[j]:end[j]] = 1000 * x[start[j]:end[j], 2] ./ V1[j];
    ceHat[start[j]:end[j]] = 1000 * x[start[j]:end[j], 4] ./ V1[j];
    respHat[start[j]:end[j]] = 100 * ceHat[start[j]:end[j]] ./
       (EC50[j] + ceHat[start[j]:end[j]]);
  }

  for(i in 1:nObs){
    cHatObs[i] = cHat[iObs[i]];
    respHatObs[i] = respHat[iObs[i]];
  }
}
                    .
                    .
                    .
  }
```

The model is simultaneously fitted to the PK and the PD data. For this effect compartment model, we construct a constant rate matrix and use `linOdeModel`. Correct use of Torsten requires the user pass the entire event history (observation and dosing events) for an individual to the function. Thus the Stan model shows the call to `linOdeModel` within a loop over the individual subjects rather than over the individual observations.

Results. We use the same diagnosis tools as for the previous example. The MCMC history plots (figure 10) suggest the 4 chains have converged to common distributions. We note some minor auto-correlations for lp_- (the log posterior) and for IIV parameters: specifically Ω_{ke_0} and ρ. The correlation matrix ρ does not explicitly appear in the model, but it is used to construct Ω, which parametrizes the PK IIV. The fits to the plasma concentration (figure 12) are in close agreement with the data, notably for the sparse data case (phase IIa study). The fits to the PD data (figure 13) look good, though the data is more noisy. The model reflects the noise by producing larger credible intervals. The estimated values of the parameters are consistent with the values used to simulate the data.

TABLE 3. Summary of the MCMC simulations of the marginal posterior distributions of the model parameters for example 2

|	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
CLHat	10.095	0.003	0.201	9.712	9.958	10.096	10.231	10.483	4000.000	0.999
QHat	14.867	0.014	0.357	14.182	14.620	14.862	15.106	15.563	678.208	1.007
V1Hat	34.188	0.067	1.089	31.940	33.494	34.214	34.918	36.251	267.748	1.016
V2Hat	103.562	0.076	2.925	98.031	101.600	103.455	105.472	109.583	488.296	1.001
kaHat	1.930	0.004	0.077	1.771	1.880	1.933	1.982	2.076	334.888	1.014
ke0Hat	1.050	0.001	0.044	0.967	1.020	1.051	1.078	1.137	164.741	1.000
EC50Hat	104.337	0.040	2.100	100.169	102.909	104.345	105.768	108.351	744.041	1.000
sigma	0.099	0.000	0.002	0.095	0.097	0.099	0.100	0.103	906.342	1.002
sigmaResp	10.156	0.003	0.197	9.779	10.023	10.154	10.286	10.552	4000.000	1.000
omega[1]	0.270	0.000	0.016	0.241	0.259	0.269	0.280	0.302	4000.000	1.001
omega[2]	0.231	0.001	0.021	0.192	0.217	0.230	0.245	0.275	531.512	1.006
omega[3]	0.219	0.002	0.031	0.158	0.199	0.218	0.238	0.281	158.198	1.017
omega[4]	0.267	0.001	0.026	0.218	0.249	0.266	0.284	0.319	684.870	1.001
omega[5]	0.285	0.002	0.037	0.214	0.259	0.284	0.309	0.361	284.545	1.009
omegaKe0	0.271	0.003	0.047	0.183	0.239	0.271	0.303	0.363	217.350	1.007
omegaEC50	0.213	0.001	0.021	0.174	0.199	0.213	0.227	0.255	190.193	1.000
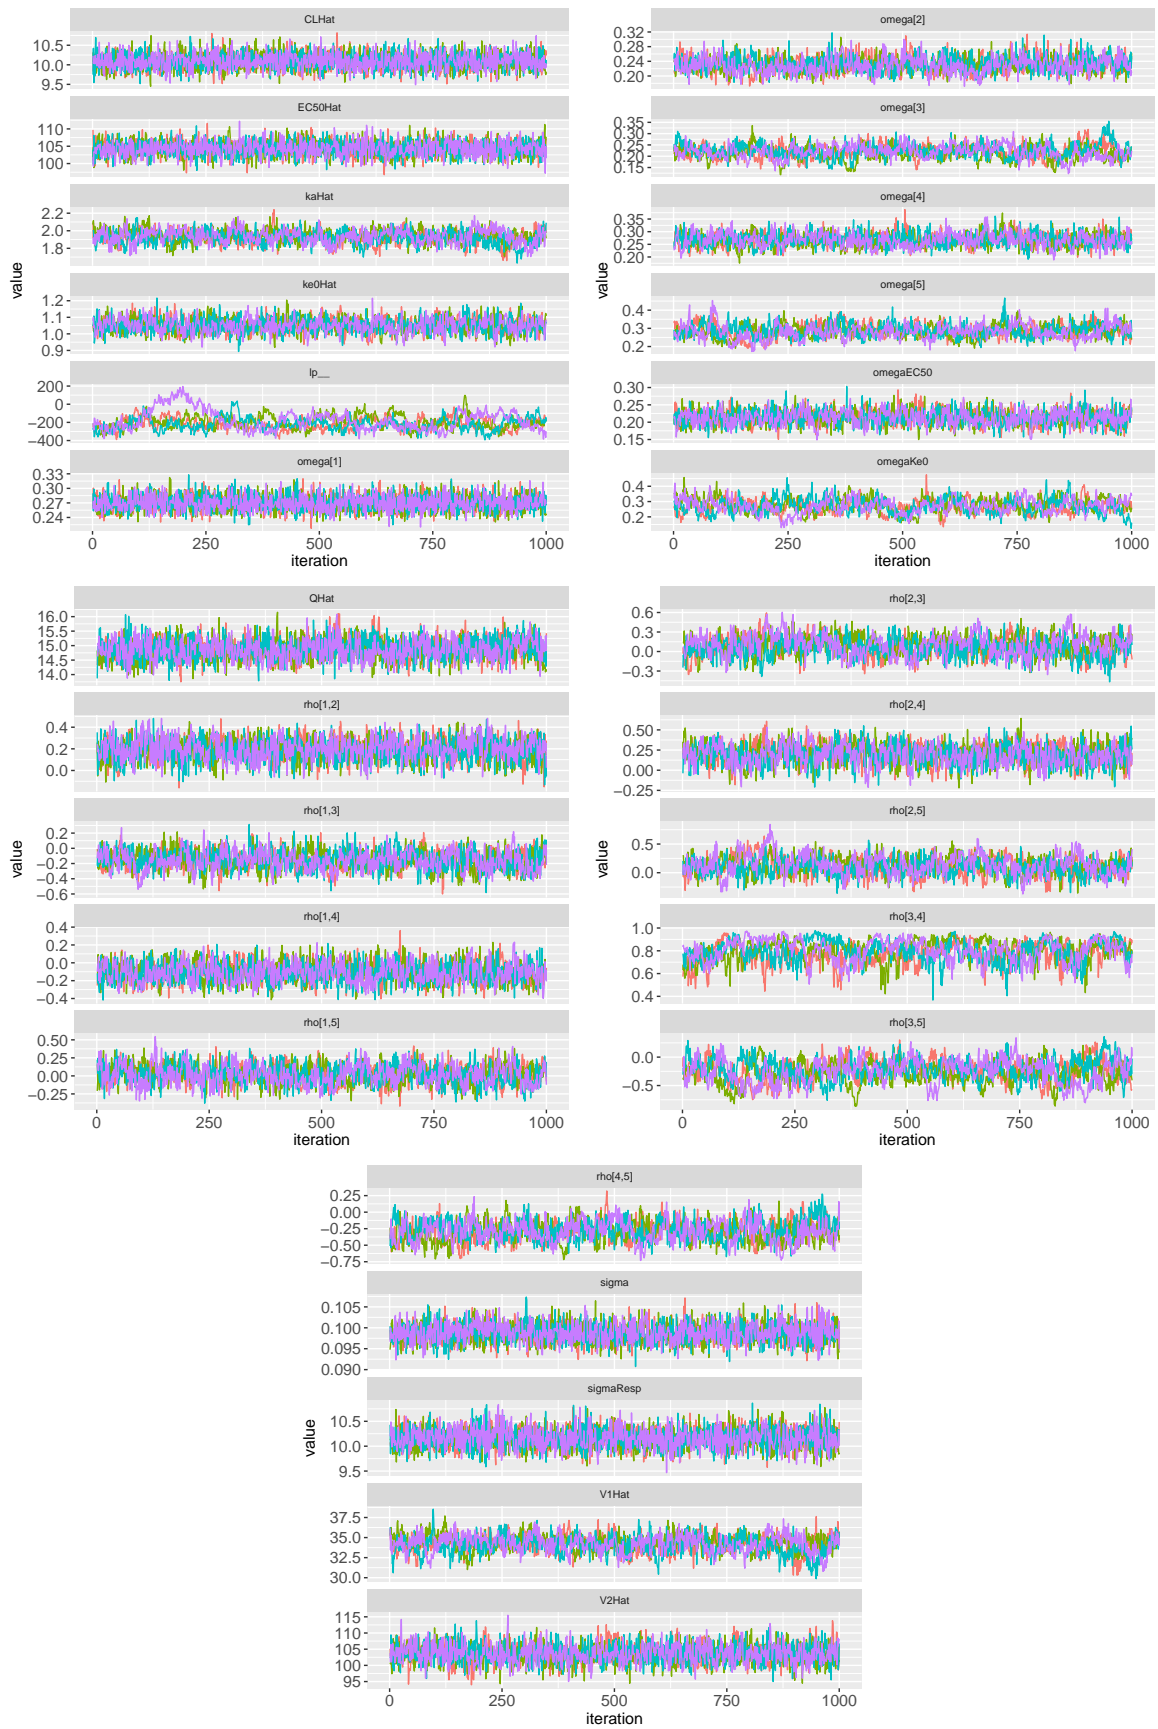
FIGURE 10. MCMC history plots for the parameters of an Effect Compartment Model (each color corresponds to a different chain) for example 2

FIGURE 11. Posterior Marginal Densities of the Model Parameters of an Effect Compartment Model (each color corresponds to a different chain) for example 2
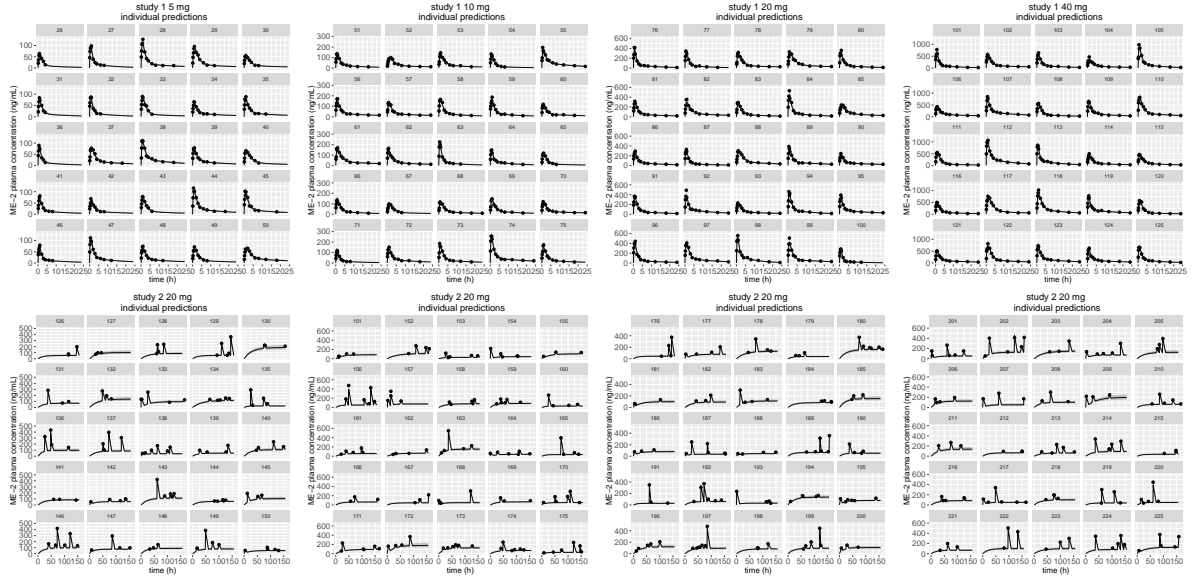

FIGURE 12. Predicted (posterior median and 90 % credible intervals) and observed plasma drug concentrations for example 2 for an Effect Compartment Model

FIGURE 13. Predicted (posterior median and 90 % credible intervals) and observed PD Response for example 2

3.2. **Friberg-Karlsson Semi-Mechanistic Model** [5]**.**

In this third example, we deal with a more sophisticated PD effect, described by a system of nonlinear ODEs. The PK effects are still described by a two compartment model with a first-order absorption.

Neutropenia is observed in patients receiving an ME-2 drug. Our goal is to model the relation between neutrophil counts and drug exposure. Using a feedback mechanism, the body maintains the number of neutrophils at a baseline value (figure 14). While in the patient's blood, the drug impedes the production of neutrophils. As a result, the neutrophil count goes down, and after the drug clears out, the feedback mechanism kicks in and brings the neutrophil count back to baseline.

Friberg-Karlsson Model for drug-induced myelosuppression (ANC)

$$
\begin{aligned}
\log(ANC_{ij}) &\sim N(Circ_{ij}, \sigma_{ANC}^2) \\
\log\left(MTT_j, Circ_{0j}, \alpha_j\right) &\sim N\left(\log\left(\widehat{MTT}, \widehat{Circ_0}, \widehat{\alpha}\right), \Omega_{ANC}\right) \\
\left(\widehat{MTT}, \widehat{Circ_0}, \widehat{\alpha}, \gamma\right) &= (125, 5, 2, 0.17) \\
\Omega_{ANC} &= \begin{pmatrix} 0.2^2 & 0 & 0 \\ 0 & 0.35^2 & 0 \\ 0 & 0 & 0.2^2 \end{pmatrix}, \quad \sigma_{ANC} = 0.1 \\
\Omega_{PK} &= \begin{pmatrix} 0.25^2 & 0 & a0 & 0 & 0 \\ 0 & 0.4^2 & 0 & 0 & 0 \\ 0 & 0 & 0.25^2 & 0 & 0 \\ 0 & 0 & 0 & 0.4^2 & 0 \\ 0 & 0 & 0 & 0 & 0.25^2 \end{pmatrix}
\end{aligned}
$$

FIGURE 14. Friberg-Karlsson semi-mechanistic Model [5]

The PK and the PD data are simulated using the following treatment.

- Phase IIa trial in patients
  - Multiple doses: 80,000 mg
  - Parallel dose escalation design
  - 15 subjects
  - PK: plasma concentration of parent drug (c)
  - PD response: Neutrophil count (ANC)
  - PK measured at 0.083, 0.167, 0.25, 0.5, 0.75, 1, 2, 3, 4, 6, 8, 12, 18, and 24 hours
  - PD measured once every two days for 28 days.

FIGURE 15. Stan language for coding an ODE system describing a Friberg-Karlsson Mechanism

```
        real[] twoCptNeutModelODE(real t,
                    real[] x,
                    real[] parms,
                    real[] rdummy,
                    int[] idummy){
    real CL = parms[1];
    real Q = parms[2];
    real V2 = parms[3];
    real V3 = parms[4];
    real ka = parms[5];
    real mtt = parms[6];
    real circ0 = parms[7];
    real gamma = parms[8];
    real alpha = parns[9];
    real k10 = CL / V2;
    real k12 = Q / V2;
    real k21 = Q / V3;
    real ktr = 4 / mtt;
    real dxdt[8];
    real conc;
    real EDrug;
    real transit1;
    real transit2;
    real transit3;
    real circ;
    real prol;

    dxdt[1] = -ka * x[1];
    dxdt[2] = ka * x[1] - (k10 + k12) * x[2] + k21 * x[3];
    dxdt[3] = k12 * x[2] - k21 * x[3];
    conc = x[2]/V1;
    EDrug = alpha * conc;
    // x[4], x[5], x[6], x[7] and x[8] are differences from circ0.
    prol = x[4] + circ0;
    transit1 = x[5] + circ0;
    transit2 = x[6] + circ0;
    transit3 = x[7] + circ0;
    circ = fmax(machine_precision(), x[8] + circ0); // Device for implementing a modeled
                                        // initial condition
    dxdt[4] = ktr * prol * ((1 - EDrug) * ((circ0 / circ)^gamma) - 1);
    dxdt[5] = ktr * (prol - transit1);
    dxdt[6] = ktr * (transit1 - transit2);
    dxdt[7] = ktr * (transit2 - transit3);
    dxdt[8] = ktr * (transit3 - circ);

    return dxdt;
}
```

Once again, we simultaneously fit the model to the PK and the PD data. From a computational perspective, this is a much more difficult problem than the one we dealt with in previous examples. The nonlinear nature of the ODEs forces us to use a numerical solver, which is significantly slower than the linear methods we have employed so far. Because the ODE system of interest is non-stiff, we use the *rk45* version of `genOdeModel`.

It pays off to construct informative priors. For instance, we could fit the PK data first, as was done in example 1, and get informative priors on the PK parameters. The PD parameters are drug independent, so we can use information from the neutropenia literature. In this example, we choose to use weakly informative priors on the PK parameters and strongly informative priors on the PD parameters.

Since it takes a long time to run the model, we only use 100 iterations per chain, and study what we can learn from this less than optimal scenario. It is worth noting that Stan, because of its highly efficient MCMC sampler, still does a reasonable job estimating the posterior distribution.

FIGURE 16. Stan language for fitting a Friberg-Karlsson model using genCptModel_rk45 (abstract)

```
transformed parameters {
                  .
                  .
                  .
  for(i in 1:nSubjects) {

    parms[1] = thetaM[i, 1] * (weight[i] / 70)^0.75; # CL
    parms[2] = thetaM[i, 2] * (weight[i] / 70)^0.75; # Q
    parms[3] = thetaM[i, 3] * (weight[i] / 70); # V1
    parms[4] = thetaM[i, 4] * (weight[i] / 70); # V2
    parms[5] = kaHat; # ka
    parms[6] = thetaM[i, 5]; # mtt
    parms[7] = thetaM[i, 6]; # circ0
    parms[8] = gamma;
    parms[9] = thetaM[i, 7]; # alpha

    x[start[i]:end[i]] = generalCptModel_rk45(twoCptNeutModelODE, 8,
                                     time[start[i]:end[i]],
                                     amt[start[i]:end[i]],
                                     rate[start[i]:end[i]],
                                     ii[start[i]:end[i]],
                                     evid[start[i]:end[i]],
                                     cmt[start[i]:end[i]],
                                     addl[start[i]:end[i]],
                                     ss[start[i]:end[i]],
                                     parms, biovar, tlag,
                                     1e-6, 1e-6, 1e8);

    cHat[start[i]:end[i]] = x[start[i]:end[i], 2] / parms[1][3]; # divide by V1
    neutHat[start[i]:end[i]] = x[start[i]:end[i], 8] + parms[1][7]; # Add baseline

  }

  for(i in 1:nObsPK) cHatObs[i] = cHat[iObsPK[i]];
  for(i in 1:nObsPD) neutHatObs[i] = neutHat[iObsPD[i]];

                  .
                  .
                  .
```

Results. The MCMC history plots are not as convincing as in the previous examples, mostly because the number of iterations is small (100 versus 1000 in the previous example). It does however look as though the chains are converging to a common distribution, and we see little auto-correlation (in particular, we expect that if we had run the model for 1000 iterations, we would obtain the desired "fuzzy caterpillar" look). The plots of the marginal posterior distributions clearly show that the chains have not (yet) converged to a common distribution, but they do not disagree significantly. Still, the need for more iterations is evident. The model fits the data, and the credible interval reflect the noise in the data. The parameters estimation reflects the real value of the parameters.

TABLE 4. Summary of the MCMC simulations of the marginal posterior distributions of the model parameters for example 3

|	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
CL	9.986	0.009	0.174	9.641	9.872	9.982	10.107	10.331	400.000	0.997
Q	14.633	0.055	1.106	12.505	13.992	14.623	15.296	16.948	400.000	0.996
V1	32.909	0.174	2.439	28.203	31.186	32.836	34.762	37.750	195.828	1.008
V2	106.631	0.311	6.226	95.234	102.269	106.403	111.000	118.533	400.000	0.999
ka	1.882	0.012	0.175	1.582	1.756	1.871	2.006	2.223	196.052	1.007
sigma	0.106	0.001	0.010	0.089	0.098	0.105	0.112	0.132	259.693	1.009
alpha	3.3E-04	1.4E-06	2.2E-05	2.9E-04	3.2E-04	3.3E-04	3.5E-04	3.8E-04	247	1.01
mtt	132.763	0.515	6.498	120.843	128.082	132.223	136.694	146.845	159.372	1.024
circ0	5.014	0.009	0.172	4.711	4.888	5.000	5.138	5.334	400.000	1.000
gamma	0.190	0.002	0.022	0.153	0.175	0.187	0.202	0.239	139.485	1.025
sigmaNeut	0.092	0.001	0.014	0.068	0.082	0.090	0.100	0.125	161.199	1.010
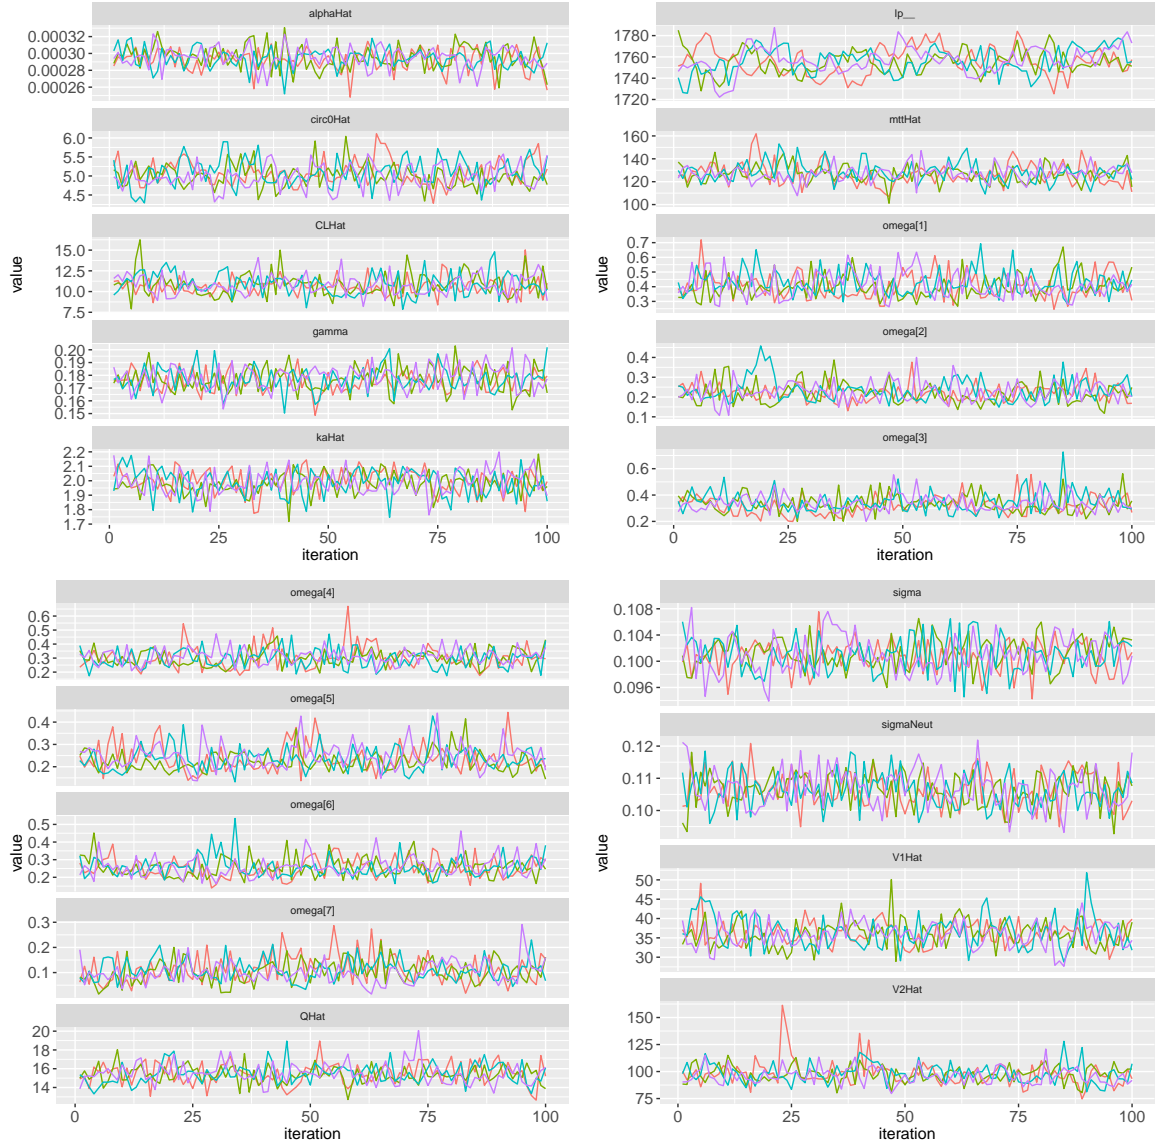
FIGURE 17. MCMC history plots for the parameters of a Friberg-Karlsson semi-mechanistic model (each color corresponds to a different chain) for example 3
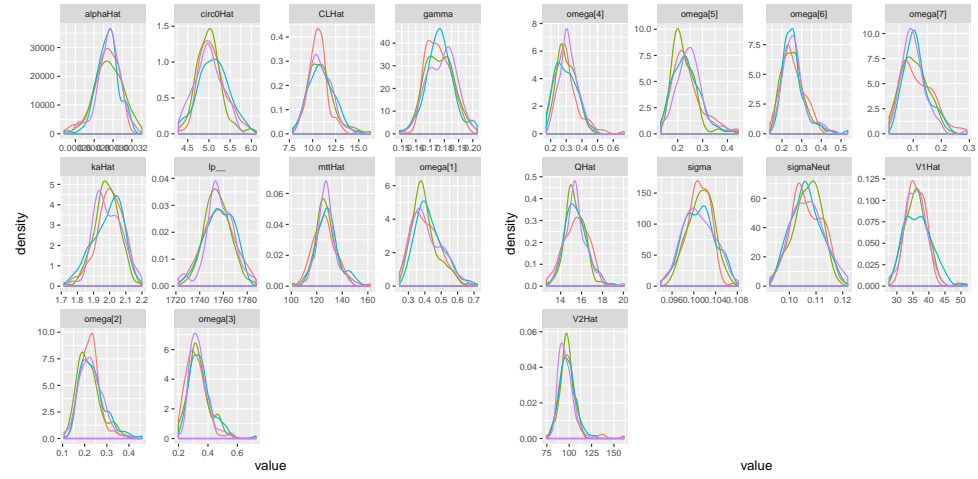
FIGURE 18. Posterior Marginal Densities of the Model Parameters of a Friberg-Karlsson semi-mechanistic model (each color corresponds to a different chain)
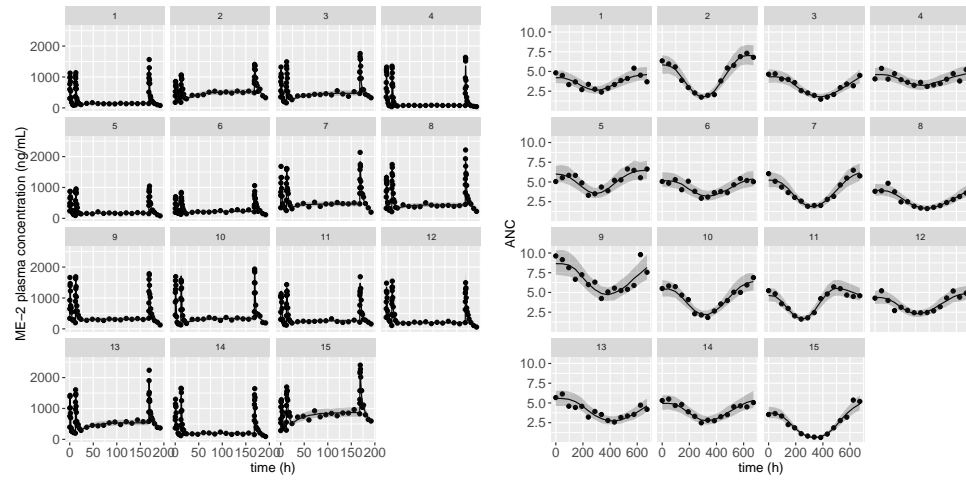

FIGURE 19. Predicted (posterior median and 90 % credible intervals) and observed plasma drug concentrations, and Neutrophil counts, for a Friberg-Karlsson semi-mechanistic model

4. UNDER THE HOOD DESIGN

We here discuss some background theory and the design of Torsten at a C++ level.

Our approach is heavily based on the *BUGS model library*, a prototype PKPD model library for WinBUGS (`https://bitbucket.org/metrumrg/bugsmodellibrary/wiki/Home`), developed by Metrum Research Group in 2009.

This section is intended for developers. No knowledge of C++ is required for users.

4.1. **Computing Amounts in an ODE-based model.**

Most PKPD model are based on ODEs that describe how PK and PD amounts evolve over time. For instance, the following ODEs describe drug diffusion in a one compartment model with a first-order absorption from the gut:

$$
\begin{aligned}
\frac{dGUT}{dt} &= -kaGUT \\
\frac{dCENT}{dt} &= kaGUT - \frac{CL}{V_1}CENT
\end{aligned}
$$

If the ODEs fully describe the PKPD system, knowing the state y_0 at time t_0 fully defines the solution at finite times. Exploiting this property, Torsten calculates the evolution of amounts in each compartment from one event to the other. The initial conditions of the ODE system are specified by the previous event, and the ODEs are integrated from $t_{previous}$ to $t_{current}$.

Note we cannot simply integrate from t_{first} to t_{last} because the ODEs do not describe exterior interventions, such as additional dosing. Torsten treats these interventions independently. Most importantly, Torsten only integrates between t_0 and t_1 if no exterior interventions occur during this interval. To achieve this, it is key to properly handle the *event schedule.*

All five functions in Torsten call the C++ function `pred`, which:

(1) augments the event schedule to include all events that alter the system
(2) calculates the amounts in each compartment at each event of the augmented schedule by:
   • computing the *natural* evolution of the system by integrating ODEs,
   • or computing alterations due to exterior interventions
(3) returns the amounts at each event of the original schedule

The Event Schedule depends on the user's input (TIME, EVID, CMT, AMT, RATE, ADDL, II, SS). The event schedule may need to be augmented if, for example, an event specifies a patient receives multiple doses at a regular time interval. Consider:

`TIME = 0, EVID = 1, CMT = 1, AMT = 1500, RATE = 0, ADDL = 4, II = 10, SS = 0`

This Event specifies that a time 0 (TIME = 0), a patient receives a 1500 mg (AMT = 1500) drug dose (EVID = 1) in the gut (CMT = 1), and will receive an additional dose every 10 hours (II = 10) until the patient has received a total of 5 doses (ADDL = 4, being the number of additional doses, + 1, the original dose). Such an Event really corresponds to 5 dosing events.

To integrate the ODEs, `pred` calls `pred1` (prediction for one event) or `predSS` (prediction for one event if the system is in a steady state, i.e SS = 1). `pred1` and `predSS` are functors and get constructed differently by each Torsten function. Under `PKModelOneCpt`, `pred1` analytically computes the solution, while under `generalCptModel_rk45` it numerically solves the ODEs.

4.2. **Structure of a Torsten Function.**

Under the structural scheme described above, a Torsten function performs a very simple set of actions:

(1) Consistent with Stan practices, check the validity of the arguments and of the parameter values
(2) Construct the PKModel object, which contains basic information about the model, such as the number of compartments
(3) Construct the pred1 and predSS functions
(4) Call pred

4.3. **Implementing Torsten in Stan.**

Modifications in Stan-math. All five Torsten functions are located under the `Torsten` directory, under `stan/math`. We modified `rev/math` to include the `torsten/torsten.hpp` header file. The code can be found on GitHub: `https://github.com/charlesm93/math`

Modifications in Stan. Further modification are done in Stan to expose the Torsten functions to the Stan language. We edited `function_signatures.h` to expose `PKModelOneCpt`, `PKModelTwoCpt`, and `linOdeModel`. The general ODE model functions are higher-order functions (i.e. they take another function as one of their arguments). They were exposed by directly modifying the grammar files, following very closely the example of `integrate_ode_rk45` and `integrate_ode_bdf`.

The code can be found on GitHub: `https://github.com/charlesm93/stan`

References

[1] Carpenter, B., Gelman, A., Hoffman, M., Lee, D., Goodrich, B., Betancourt, M., Brubaker, M.A., Guo, J., Li, P. and Riddel, A. Stan: A probablistic programming language. *Journal of Statistical Software (in press)* (2016).

[2] Hoffman, M.D. and Gelman, A. The no-u-turn sampler: Adaptively setting path lengths in hamiltonian monte carlo. *Journal of Machine Learning Research* (2014):1593–1623.

[3] Carpenter, B., Hoffman, M.D., Brubaker, M.A., Lee, D., Li, P. and Betancourt, M.J. The stan math library: Reverse-mode automatic differentiation in c++. *arXiv 1509.07164.* (2015).

[4] Moler, C. and Van Loan, C. Nineteen dubious ways to compute the exponential of a matrix, twenty-five years later. *SIAM Review* (2003).

[5] Friberg, L.E. and Karlsson, M.O. Mechanistic models for myelosuppression. *Invest New Drugs* **21** (2003):183–194.